

LA03_Ex2_KDE

April 28, 2018

1 Team Members

1.1 RaviKiran Bhat

1.2 Rubanraj Ravichandran

1.3 Mohammad Wasil

1.4 Ramesh Kumar

2 Task1

2.1 Compare the outcomes of different implementations of KDEs.

There are several options available for computing KDE in Python. - SciPy: gaussian_kde. - Statsmodels: KDEUnivariate and KDEMultivariate. - Scikit-learn: KernelDensity.

2.2 1). Generate synthetic data and plot them

Generate synthetic dataset the distribution of which can be presented as a combination of three Gaussian distributions with the following parameters: $\mu_1=1$, $\sigma_1=1$ and $\mu_2=8$, $\sigma_2=2$ and $\mu_3=14$, $\sigma_3=1.5$. Generate 1000 samples from the distribution. Plot the pdf of this distribution and the generated samples. 3) Use the generated samples to perform - (i) KDE with Scipy, - (ii) Univariate KDE with Statsmodels, - (iii) Multivariate KDE with Statsmodels as well as - (iv) KDE with Scikit-learn. 4) Plot all four distributions on one figure.

```
In [1]: import numpy as np
        from scipy.stats import norm
        import pandas as pd
        import matplotlib.pyplot as plt
        from scipy import stats
        from statsmodels.nonparametric.kde import KDEUnivariate
        from statsmodels.nonparametric.kernel_density import KDEMultivariate
        from sklearn.neighbors import KernelDensity
```

```
%matplotlib inline
```

```
In [2]: np.random.seed(0)
        gaussian1 = 1 + 1 * np.random.randn(1000)
```

```

gaussian2 = 8 + 2 * np.random.randn(1000)
gaussian3 = 14 + 1.5 * np.random.randn(1000)
gaussian_mixture = np.hstack([gaussian1, gaussian2, gaussian3])

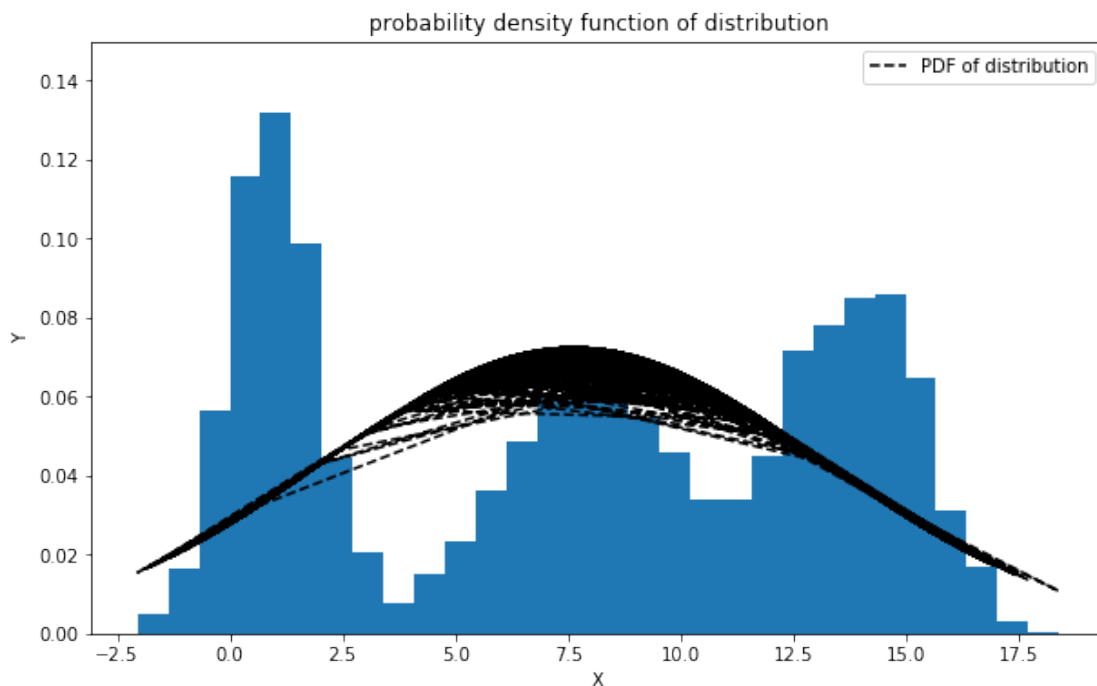
df = pd.DataFrame(gaussian_mixture, columns=['data'])
# parametric fit: assume normal distribution
param_density = stats.norm.pdf(gaussian_mixture, np.mean(gaussian_mixture), np.std(gaussian_mixture))

fig, ax = plt.subplots(figsize=(10, 6))
ax.hist(df.values, bins=30, normed=True)
ax.plot(df, param_density, 'k--', label='PDF of distribution')
ax.set_ylim([0, 0.15])
ax.set_xlabel("X")
ax.set_ylabel("Y")
ax.set_title("probability density function of distribution")
ax.legend(loc='best')

```

/home/ramesh/anaconda2/lib/python2.7/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: T
 warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[2]: <matplotlib.legend.Legend at 0x7f3921c5f0d0>



```

In [3]: # non-parametric pdf
fig, ax = plt.subplots(figsize=(10, 6))

```

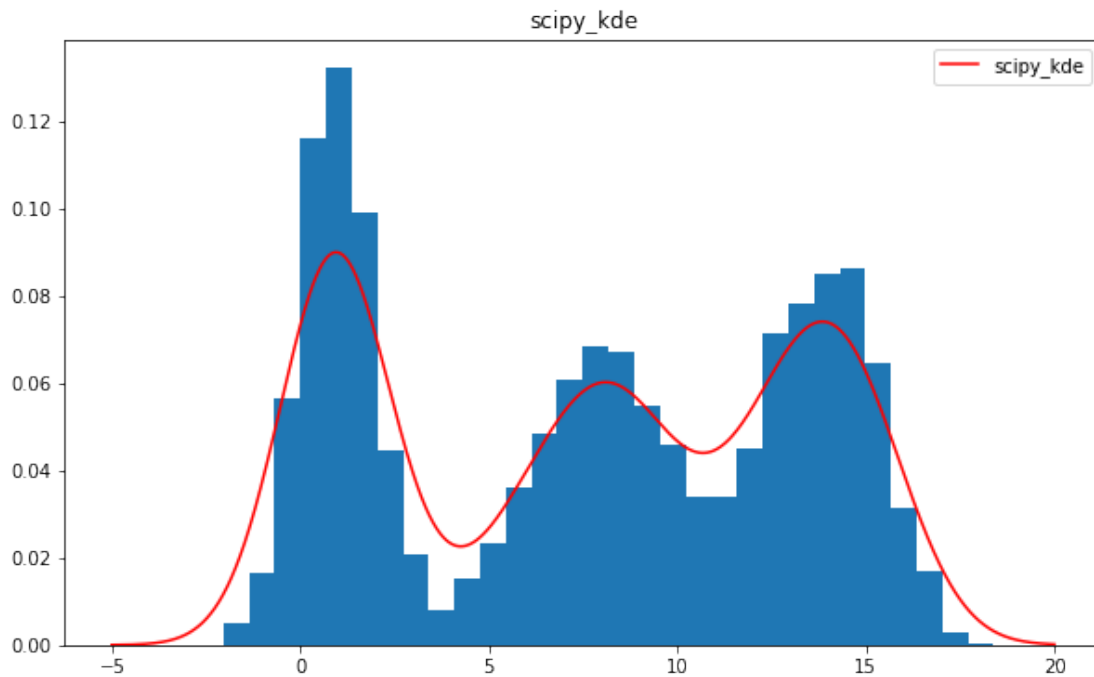
```

ax.hist(df.values, bins=30, normed=True)

scipy_kde = stats.gaussian_kde(df.values.ravel())
x = np.linspace(-5, 20, 3000)
scipy_kde = scipy_kde(x)
ax.plot(x, scipy_kde, 'r-', label='scipy_kde')
ax.legend(loc='best')
plt.title('scipy_kde')

```

Out[3]: Text(0.5,1,'scipy_kde')



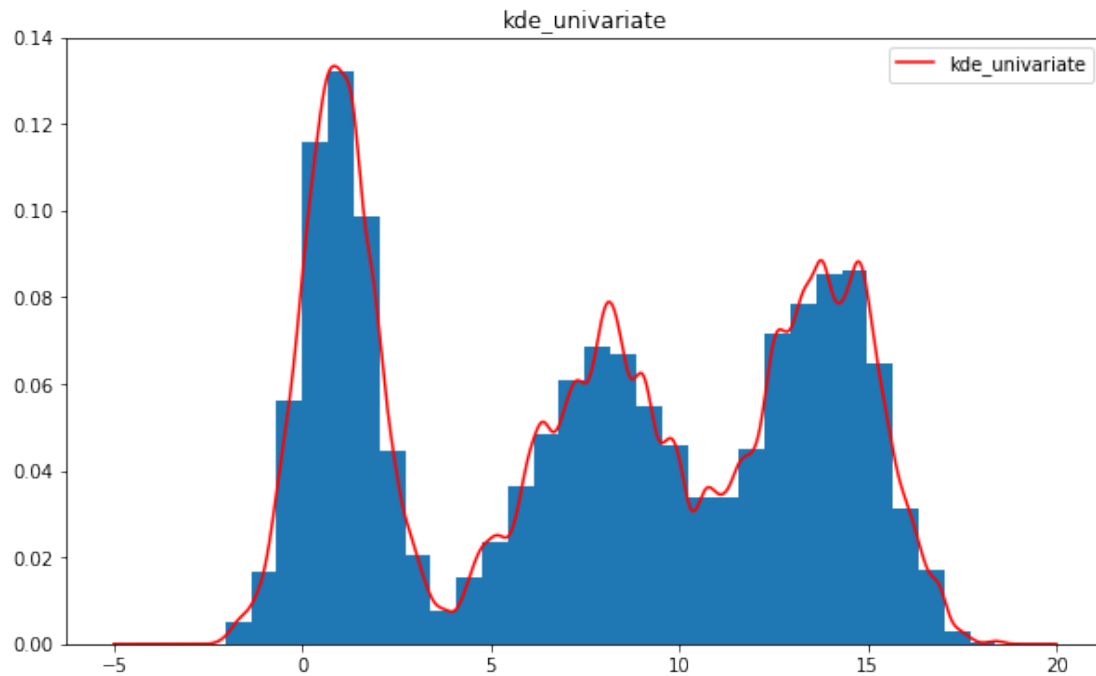
```

In [4]: fig, ax = plt.subplots(figsize=(10, 6))
ax.hist(df.values, bins=30, normed=True)

kde_univariate = KDEUnivariate(df.values.ravel())
kde_univariate.fit(bw=0.2)
x = np.linspace(-5, 20, 3000)
pdf = kde_univariate.evaluate(x)
ax.plot(x, pdf, color='red', label='kde_univariate')
ax.legend(loc='best')
plt.title('kde_univariate')

```

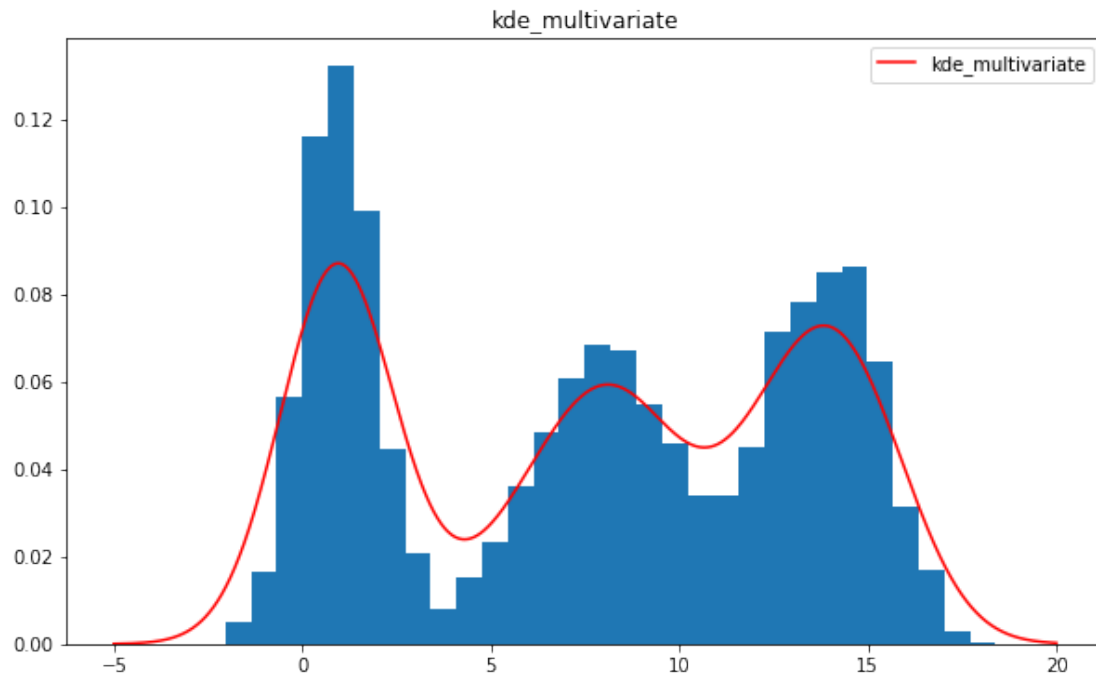
Out[4]: Text(0.5,1,'kde_univariate')



```
In [5]: fig, ax = plt.subplots(figsize=(10, 6))
        ax.hist(df.values, bins=30, normed=True)

        kde = KDEMultivariate(df.values, var_type='c')
        x_grid = np.linspace(-5, 20, 3000)
        kde_multivariate = kde.pdf(x_grid)
        plt.plot(x_grid, kde_multivariate, color='red', label='kde_multivariate')
        ax.legend(loc='best')
        plt.title('kde_multivariate')

Out[5]: Text(0.5,1,'kde_multivariate')
```

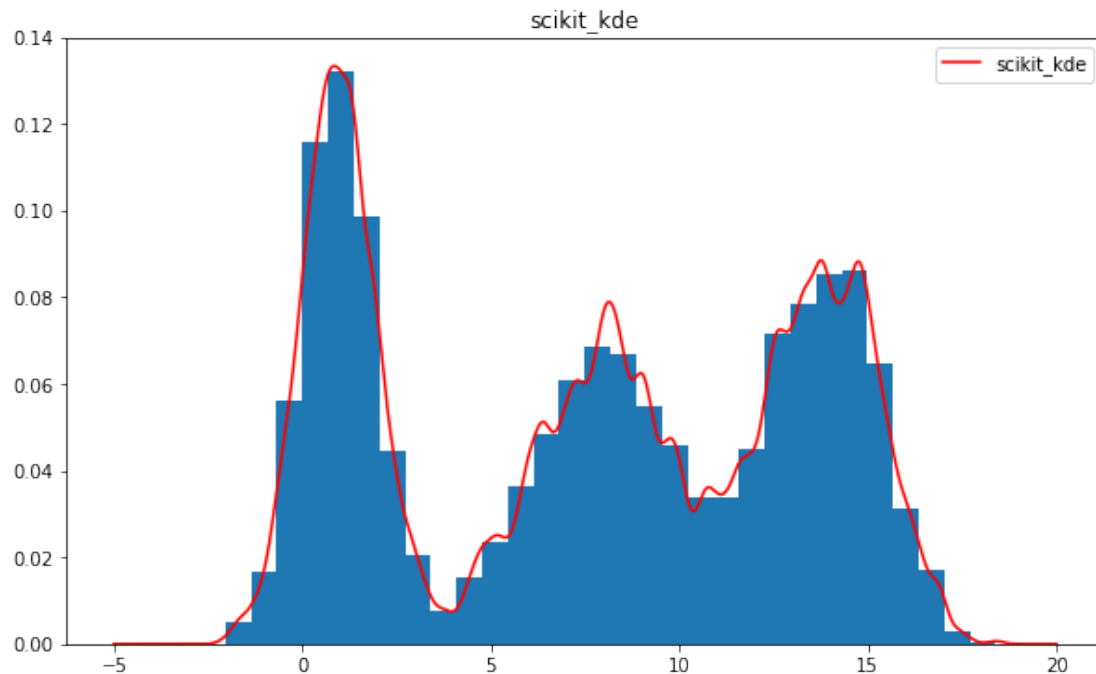


In [6]: # KDE Scikit learn

```
fig, ax = plt.subplots(figsize=(10, 6))
ax.hist(df.values, bins=30, normed=True)

scikit_kde = KernelDensity(kernel='gaussian', bandwidth=0.2).fit(df.values.ravel(), np
x = np.linspace(-5, 20, 3000)
log_dens = scikit_kde.score_samples(x[:, np.newaxis])
ax.plot(x, np.exp(log_dens), label="scikit_kde", color='red')
ax.legend(loc='best')
plt.title('scikit_kde')
```

Out[6]: Text(0.5,1,'scikit_kde')



```
In [10]: # Plot all four distributions on one figure.
fig, ax = plt.subplots(figsize=(10, 6))
ax.hist(df.values, bins=30, normed=True)
ax.plot(x, scipy_kde, 'r-', label='scipy_kde')

ax.plot(x, pdf, color='green', label='kde_univariate')

ax.plot(x, kde_multivariate, color='yellow', label='kde_univariate')

ax.plot(x, np.exp(log_dens), label="scikit_kde", color='black')
ax.legend(loc='best')

ax.set_title("Comparison between all four distributions")
ax.set_xlabel("x")
ax.set_ylabel("y")
```

```
Out[10]: Text(0,0.5,'y')
```

