

as07

June 10, 2018

1 Team Members

1.1 RaviKiran Bhat

1.2 Rubanraj Ravichandran

1.3 Mohammad Wasil

1.4 Ramesh Kumar

```
In [1]: import tensorflow as tf
        from sklearn.linear_model import LinearRegression, LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.metrics import log_loss
        import numpy as np
        import pickle
        from sklearn.externals import joblib
        import matplotlib.pyplot as plt
        %matplotlib inline
```

```
/home/ramesh/anaconda2/lib/python2.7/site-packages/h5py/__init__.py:34: FutureWarning: Conversion
from ._conv import register_converters as _register_converters
```

```
In [2]: #Load MNIST data from tf
        from tensorflow.examples.tutorials.mnist import input_data
        mnist = input_data.read_data_sets("MNIST_data/", one_hot=False)
        mnist_one_hot = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

```
WARNING:tensorflow:From <ipython-input-2-64bc5dd34d72>:3: read_data_sets (from tensorflow.contrib.learn.python.learn.python_lib) is deprecated and will be removed in a future version.
Instructions for updating:
```

```
Please use alternatives such as official/mnist/dataset.py from tensorflow/models.
```

```
WARNING:tensorflow:From /home/ramesh/anaconda2/lib/python2.7/site-packages/tensorflow/contrib/learn/python/learn/python_lib.py:114: read_data_sets (from tensorflow.contrib.learn.python.learn.python_lib) is deprecated and will be removed in a future version.
Instructions for updating:
```

```
Please write your own downloading logic.
```

```
WARNING:tensorflow:From /home/ramesh/anaconda2/lib/python2.7/site-packages/tensorflow/contrib/learn/python/learn/python_lib.py:114: read_data_sets (from tensorflow.contrib.learn.python.learn.python_lib) is deprecated and will be removed in a future version.
Instructions for updating:
```

```
Please use tf.data to implement this functionality.
```

```
Extracting MNIST_data/train-images-idx3-ubyte.gz
```

```

WARNING:tensorflow:From /home/ramesh/anaconda2/lib/python2.7/site-packages/tensorflow/contrib/le
Instructions for updating:
Please use tf.data to implement this functionality.
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
WARNING:tensorflow:From /home/ramesh/anaconda2/lib/python2.7/site-packages/tensorflow/contrib/le
Instructions for updating:
Please use alternatives such as official/mnist/dataset.py from tensorflow/models.
Extracting MNIST_data/train-images-idx3-ubyte.gz
Extracting MNIST_data/train-labels-idx1-ubyte.gz
WARNING:tensorflow:From /home/ramesh/anaconda2/lib/python2.7/site-packages/tensorflow/contrib/le
Instructions for updating:
Please use tf.one_hot on tensors.
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz

```

```

In [3]: def plot_img(image, index):
        f, ax = plt.subplots(1, len(index))
        for i in range(len(index)):
            ax[i].imshow(np.reshape(image[index[i]], (28,28)), cmap='Greys')
            ax[i].set_yticklabels([])
            ax[i].set_xticklabels([])
        plt.show()

```

```

In [24]: class mnistTwoClassifiers(object):
        def __init__(self):
            self.step = 100

        def update_datasets(self, train_images, train_labels, test_images, test_labels):
            self.train_images = train_images
            self.train_labels = train_labels
            self.test_images = test_images
            self.test_labels = test_labels

        def logistic_regression(self, train_size):
            self.logistic_model = LogisticRegression()
            print self.train_images[:train_size].shape
            print self.train_labels[:train_size].shape
            self.logistic_model.fit(self.train_images[:train_size], self.train_labels[:train_size])
            return self.logistic_model

        def random_forest(self, train_size):
            self.random_forest_model = RandomForestClassifier()
            %time self.random_forest_model.fit(self.train_images[:train_size], self.train_labels[:train_size])
            return self.random_forest_model

```

```

def generate_adversarial_example(self, num_of_samples, clf, epsilon, test_images, t
    adv_samples = np.zeros((num_of_samples, 784))
    adv_labels = np.zeros((num_of_samples, 10))
    adv_labels_true = np.zeros((num_of_samples, 1))
    for i in range(num_of_samples):
        y_predict = None
        if clf == "logistic_regression":
            y_predict = self.logistic_model.predict_proba(test_images[i:i+1])
        elif clf == "random_forest":
            y_predict = self.random_forest_model.predict_proba(test_images[i:i+1])

        y_true = test_labels[i:i+1]
        y_true_index = np.where(y_true == 1)[1][0]
        predictions = y_predict
        error = (predictions - y_true)**(2)
        error = error[0][y_true_index]
        gradient = error * train_images[i:i+1]
        gradient /= len(train_images[i:i+1])
        signs = np.sign(gradient)
        img_adversarial = epsilon * signs * test_images[i:i+1]
        adv_samples[i] = img_adversarial
        adv_labels[i] = y_true
        adv_labels_true[i] = y_true_index
    adv_samples = np.asarray(adv_samples)
    return adv_samples, adv_labels, adv_labels_true

def test(self, test_size, classifier):
    if classifier == "logistic_regression":
        print(self.logistic_model.score(self.test_images[:test_size], self.test_lab
    elif classifier == "random_forest":
        print(self.random_forest_model.score(self.test_images[:test_size], self.tes

def predict(self, model, test_images, test_size):
    return model.predict(test_images[:test_size])

def test_with_adversarial(self, model, adv_example):
    return model.predict(adv_example)

```

```

In [25]: train_images = mnist.train.images
        train_labels = mnist.train.labels
        test_images = mnist.test.images
        test_labels = mnist.test.labels
        test_labels_one_hot = mnist_one_hot.test.labels
        test_images_one_hot = mnist_one_hot.test.images
        train_labels_one_hot = mnist_one_hot.train.labels
        train_images_one_hot = mnist_one_hot.train.images
        two_clfs = mnistTwoClassifiers()

```

```
two_clfs.update_datasets(train_images, train_labels, test_images, test_labels)
```

2 Task 1 -5

```
In [26]: #train the models
logistic_model = two_clfs.logistic_regression(55000)
random_forest_model = two_clfs.random_forest(55000)
clf_logistic = two_clfs.logistic_model
clf_random_forest = two_clfs.random_forest_model
#Save classifier to pickle file
joblib.dump(clf_logistic, 'logistic_regression.pkl')
joblib.dump(clf_random_forest, 'random_forest_model.pkl')
```

```
(55000, 784)
```

```
(55000,)
```

```
CPU times: user 7.95 s, sys: 4.04 ms, total: 7.95 s
```

```
Wall time: 8.14 s
```

```
Out[26]: ['random_forest_model.pkl']
```

```
In [7]: #Load classifier
# logistic_clf = joblib.load('/tmp/logistic_regression.pkl')
# random_forest_clf = joblib.load('/tmp/random_forest_model.pkl')
```

3 Test trained Model

```
In [27]: #test the model
```

```
two_clfs.test(test_images.shape[0], "logistic_regression")
two_clfs.test(test_images.shape[0], "random_forest")
```

```
0.9198
```

```
0.9476
```

4 Use Test dataset to generate Adversarial Examples

```
In [28]: #Use test dataset
num_adv_example = 100
epsilon = 0.007
adv_imgs_logistic, labels_logistic, adv_labels_true_logistic = two_clfs.generate_adversarial_examples(
    epsilon, test_images_one_hot, test_labels)
adv_imgs_rnd_forest, labels_rnd_forest, adv_labels_true_rnd_forest = two_clfs.generate_adversarial_examples(
    epsilon, test_images_one_hot, test_labels)
```

```
In [29]: #evaluate trained model using random forest on adversarial examples
clf_random_forest.score(adv_imgs_logistic, test_labels[:num_adv_example])
```

Out[29]: 0.27

```
In [30]: #evaluate trained model using logistics on adversarial examples
        clf_logistic.score(adv_imgs_rnd_forest, train_labels[:num_adv_example])
```

Out[30]: 0.08

5 Task 6

6 Generate 55000 adversarial examples

```
In [42]: #Use original training dataset and generate adversarial examples
        num_adv_example_task_six = 55000
        epsilon_task_six = 0.007
        adv_imgs_logistic_task_six, logic_labels_task_six, logic_labels_true_task_six = two_clf
                                                    epsilon, train_images_one_hot, train_l
        adv_imgs_rnd_forest_task_six, forest_labels_task_six, forest_labels_true_task_six = two
                                                    epsilon, train_images_one_hot, train_l

In [43]: new_training_set = np.concatenate((train_images,adv_imgs_logistic_task_six))
        new_training_set_label = np.concatenate((train_labels,logic_labels_true_task_six.T[0]))
        new_training_set_label_hot = np.concatenate((train_labels_one_hot,logic_labels_task_six

In [44]: adv_classifiers = mnistTwoClassifiers()
        adv_classifiers.update_datasets(new_training_set, new_training_set_label, test_images,

In [45]: #train the models
        logistic_model = adv_classifiers.logistic_regression(new_training_set.shape[0])
        random_forest_model = adv_classifiers.random_forest(new_training_set.shape[0])
        clf_logistic = adv_classifiers.logistic_model
        clf_random_forest = adv_classifiers.random_forest_model
```

(65000, 784)

(65000,)

CPU times: user 10 s, sys: 108 ms, total: 10.1 s

Wall time: 10.5 s

7 Evaluate model trained on Adversarial examples and training dataset

```
In [49]: print "Classification accuracy on adversarial dataset using logistic regression"
        adv_classifiers.test(test_images.shape[0], "logistic_regression")
```

Classification accuracy on adversarial dataset using logistic regression
0.9199

```
In [50]: print "Classification accuracy on adversarial dataset using random forest"
        adv_classifiers.test(test_images.shape[0], "random_forest")
```

Classification accuracy on adversarial dataset using random forest
0.9465

8 Does classification performance improve?

Yes, classification performance of models trained on combination of adversarial examples and original dataset improves significantly

9 Is the new model less or more susceptible to adversarial examples?

New model is more robust with adversarial examples because it is trained on adversarial examples also.

10 Do you think you can use a regularization method in order to make the model less susceptible to adversarial examples?

No, state of the art shows that Generic regularization strategies such as dropout, pretraining, and model averaging do not confer a significant reduction in a model's vulnerability to adversarial examples, but changing to nonlinear model families such as RBF networks can do so.

```
In [51]: # np.save("adv_imgs_logistic", adv_imgs_logistic_task_six)
         # np.save("logic_labels", logic_labels_task_six)
         # np.save("adv_imgs_rnd_forest", adv_imgs_rnd_forest_task_six)
         # np.save("forest_labels", forest_labels_task_six)
```